

imd-pipeline-08march2024

March 9, 2024

The IMD Machine Learning Model: Credit to: Isaac Okello Opiro The following were performed in this step: 1. Take a peek at the raw data. 2. Review the dimensions of my dataset. 3. Review the data types of attributes in the data.

```
[202]: #Import libraries
from pycaret import*
import pandas as pd #this is for dataframe manipulation
import numpy as np #this is for numerical / mathematical computing
import matplotlib.pyplot as plt #this is for visualisation
from IPython.display import display
```

```
[203]: #import data
IMD_data = pd.read_excel ("DATASET08MARCH2024.xlsx")
```

```
[204]: #EDA
IMD_data.head(5)
```

```
[204]:
```

	patid	age	agecatak	AgeCategory	sex1	religion1	\
0	CHAKA/01/01/0001	11.668720	9 to 12	Adolescent	Female	Born Again	
1	CHAKA/01/01/0002	15.000000	13 to 17	Adolescent	Male	Protestant	
2	CHAKA/01/01/0003	6.020534	9 to 12	Adolescent	Female	Muslim	
3	CHAKA/01/02/0152	5.242984	5 to 8	Children	Female	Catholic	
4	CHAKA/01/02/0153	11.696100	9 to 12	Adolescent	Female	Muslim	

	childeduc1	heightst1	weightst1	height_m	...	stin2vntr_	httlpr1	\
0	Pre-primary	134.0	31.0	1.340	...	NaN	3.0	
1	Pre-primary	137.0	35.0	1.370	...	3.0	3.0	
2	Pre-primary	120.1	19.0	1.201	...	1.0	1.0	
3	Pre-primary	99.0	15.0	0.990	...	3.0	2.0	
4	Pre-primary	131.6	30.0	1.316	...	3.0	1.0	

	HTTLPRrs35531	rs35531	ptsd	gad	mdd	pc	intdis	IMDs
0	NaN	NaN	0.0	0.0	1.0	1	1	Yes
1	7.0	1.0	0.0	4.0	6.0	1	10	Yes
2	2.0	2.0	0.0	0.0	3.0	0	3	Yes
3	6.0	2.0	0.0	0.0	0.0	0	0	No
4	1.0	1.0	1.0	3.0	3.0	0	7	Yes

[5 rows x 45 columns]

```
[205]: # List variable names
variable_names = IMD_data.columns.tolist()

# Print the variable names
print(variable_names)
```

```
['patid', 'age', 'agecat1', 'AgeCategory', 'sex1', 'religion1', 'childeduc1',
'heightst1', 'weightst1', 'height_m', 'BMI', 'rounded_bmi', 'BMI_category',
'childtrib1', 'motherali1', 'fatherali1', 'orphanhood', 'ses', 'ses_cat',
'livelihood1', 'sexualever1', 'childstay1', 'childartk1', 'childworst1',
'childpremt1', 'chilborhiv', 'tobacco_status', 'alcohol1', 'Stress',
'GroupCategory', 'cd4takeoff1', 'CD4_category', 'Virallload',
'Virallload_Category', 'tlbase', 'stin2vntr_', 'httlpr1', 'HTTLPRrs35531',
'rs35531', 'ptsd', 'gad', 'mdd', 'pc', 'intdis', 'IMDs']
```

```
[206]: #Dropping Non-Important and Redundant Variables
# List of variables to drop
#Give DummyClass
#columns_to_drop = ['patid', 'height_m', 'motherali1', 'fatherali1',
↳ 'rounded_bmi', 'BMI', 'ses', 'ptsd', 'gad', 'mdd', 'pc', 'intdis']
#Gives best score, with AdaBoost
#columns_to_drop = ['patid', 'height_m', 'motherali1', 'fatherali1',
↳ 'rounded_bmi', 'BMI', 'ses', 'pc', 'intdis']
#Removing
columns_to_drop = ['patid', 'height_m', 'motherali1', 'fatherali1',
↳ 'rounded_bmi', 'BMI', 'ses', 'pc', 'intdis']

#columns_to_drop = ['patid', 'height_m', 'motherali1', 'fatherali1',
↳ 'rounded_bmi', 'BMI', 'ses', 'pc', 'intdis']

# Dropping the specified columns
IMD_data = IMD_data.drop(columns=columns_to_drop)
```

Note: Removing MDD, PTSD, and GAD from the data drops performance from 90s% to 70s %

Dimensions of the Data I had to check how much data I have, both in terms of rows and columns. - Too many rows and algorithms may take too long to train. Too few and perhaps I do not have enough data to train the algorithms. - Too many features and some algorithms can be distracted or super poor performance due to the curse of dimensionality.

```
[207]: print("Number of rows:", IMD_data.shape[0])
print("Number of columns:", IMD_data.shape[1])
#OR
print("The dimension is:", IMD_data.shape)
```

Number of rows: 736
Number of columns: 36
The dimension is: (736, 36)

Data Type For Each Attribute The type of each attribute is important. Strings may need to be converted to floating point values or integers to represent categorical or ordinal values.

```
[208]: IMD_data.dtypes
```

```
[208]: age                float64
agecatak                object
AgeCategory             object
sex1                    object
religion1               object
childeduc1              object
heightst1               float64
weightst1               float64
BMI_category            object
childtrib1              object
orphanhood              object
ses_cat                 object
livelihood1             object
sexualever1            object
childstay1              object
childartk1              object
childworst1             object
childpremt1             object
chilborhiv              object
tobacco_status          object
alcohol1                object
Stress                  object
GroupCategory           object
cd4takeoff1             float64
CD4_category            object
Virallload              float64
Virallload_Category     object
tlbase                  float64
stin2vntr_              float64
httlpr1                 float64
HTTLPRrs35531           float64
rs35531                 float64
ptsd                    float64
gad                     float64
mdd                     float64
IMDs                    object
dtype: object
```

Descriptive Statistics Descriptive statistics can give a great insight into the shape of each attribute.

```
[209]: IMD_data.describe().T
```

```
[209]:
```


	count	mean	std	min	25% \
age	696.0	10.574136	3.141758	4.525667	7.874744
heightst1	733.0	131.050887	22.690718	12.500000	119.000000
weightst1	736.0	29.878111	11.915348	10.000000	21.000000
cd4takeoff1	732.0	945.527322	533.852160	11.000000	568.750000
Virallload	611.0	30247.690671	155631.745971	0.000000	0.000000
tlbase	613.0	1.147457	0.362307	0.019943	0.914130
stin2vntr_	658.0	2.474164	0.643503	1.000000	2.000000
httlpr1	698.0	1.399713	0.583710	1.000000	1.000000
HTTLPRrs35531	685.0	2.486131	1.765954	1.000000	1.000000
rs35531	685.0	1.379562	0.550537	1.000000	1.000000
ptsd	728.0	0.453297	0.873271	0.000000	0.000000
gad	722.0	2.308864	2.668490	0.000000	0.000000
mdd	722.0	1.537396	2.361119	0.000000	0.000000

	50%	75%	max
age	10.368240	13.000000	1.711431e+01
heightst1	130.600000	144.000000	4.272000e+02
weightst1	27.000000	36.000000	1.530000e+02
cd4takeoff1	888.000000	1224.250000	3.924000e+03
Virallload	10.000000	801.500000	2.621863e+06
tlbase	1.133582	1.358633	2.179359e+00
stin2vntr_	3.000000	3.000000	3.000000e+00
httlpr1	1.000000	2.000000	3.000000e+00
HTTLPRrs35531	2.000000	3.000000	7.000000e+00
rs35531	1.000000	2.000000	3.000000e+00
ptsd	0.000000	1.000000	5.000000e+00
gad	2.000000	4.000000	1.800000e+01
mdd	0.000000	2.000000	1.500000e+01

Dependent Variable (DV) Note that: The DV in this cases was being diagnosed with Internalizing Mental Disorder, IMD (Yes) or Not (No), This was calculated by ensuring all participants without any form of GAD, MDD, or PTSD, was assigned as “Not having IMD”, while all those with any form of GAD, MDD, or PTSD, was amrked as having IMD.

```
[210]: ##Plotting the distribution of the dependent variable (DV).
IMD_data.groupby('IMDs').size().plot(kind='bar')
```

```
[210]: <Axes: xlabel='IMDs'>
```



```
[211]: #Checking at the distribution of the data
#IMD_data.hist()
```

NOTE: The above was not useful as the package (pycaret) could perform transformations in all datasets. Unlike in scikit-learn that I would have to decide the data to transform based on the above plot.

MODEL BUILDING USING PYCARET Since this was a classification problem (Supervised machine learning), the classification package from pycaret was used. This package contains several algorithms:

- Logistic Regression
- Support Vector Machines
- Decision Trees
- Naive Bayes
- K-Nearest Neighbors
- Random Forests

```
[212]: #Import classification model from pycaret
from pycaret.classification import *
```

```
[213]: # Setting up the data for machine learning modeling using the pycaret setup
↳function
# IMD_data: Your dataset or data structure
# target='IMDs': Specifying 'IMDs' as the target variable to predict
# session_id=123: Setting a session ID (seed) for reproducibility

s = setup(IMD_data, target='IMDs', session_id=123)
```

<pandas.io.formats.style.Styler at 0x284892b2e10>

```
[214]: #Model training
# Comparing and selecting the best-performing model based on default evaluation
↳metric
best = s.compare_models()

# Retrieve the metrics dataframe for all compared models
metrics_df = pull()
print(metrics_df)
```

<IPython.core.display.HTML object>

<pandas.io.formats.style.Styler at 0x28486640f90>

Processing: 0% | 0/61 [00:00<?, ?it/s]

	Model	Accuracy	AUC	Recall	Prec.	\
ada	Ada Boost Classifier	0.9980	0.0	0.9980	0.9981	
lightgbm	Light Gradient Boosting Machine	0.9980	0.0	0.9980	0.9981	
gbc	Gradient Boosting Classifier	0.9922	0.0	0.9922	0.9924	
qda	Quadratic Discriminant Analysis	0.9864	0.0	0.9864	0.9869	
dt	Decision Tree Classifier	0.9845	0.0	0.9845	0.9850	
rf	Random Forest Classifier	0.9805	0.0	0.9805	0.9814	
lr	Logistic Regression	0.9532	0.0	0.9532	0.9586	
et	Extra Trees Classifier	0.8639	0.0	0.8639	0.8674	
ridge	Ridge Classifier	0.8037	0.0	0.8037	0.7979	
lda	Linear Discriminant Analysis	0.8036	0.0	0.8036	0.8065	
dummy	Dummy Classifier	0.7437	0.0	0.7437	0.5531	
nb	Naive Bayes	0.7340	0.0	0.7340	0.6804	
knn	K Neighbors Classifier	0.6970	0.0	0.6970	0.6468	
svm	SVM - Linear Kernel	0.6312	0.0	0.6312	0.6213	

	F1	Kappa	MCC	TT (Sec)
ada	0.9980	0.9947	0.9948	0.271
lightgbm	0.9980	0.9947	0.9948	0.405
gbc	0.9921	0.9791	0.9796	0.333
qda	0.9862	0.9630	0.9642	0.211
dt	0.9845	0.9593	0.9600	0.188
rf	0.9801	0.9468	0.9489	0.312
lr	0.9542	0.8830	0.8870	0.304

et	0.8510	0.5893	0.6190	0.294
ridge	0.7936	0.4422	0.4557	0.186
lda	0.8013	0.4773	0.4846	0.183
dummy	0.6344	0.0000	0.0000	0.190
nb	0.6795	0.1087	0.1321	0.186
knn	0.6577	0.0498	0.0601	0.200
svm	0.5934	0.0232	0.0331	0.182

```
[215]: #Print the best model
print(best)
```

```
AdaBoostClassifier(algorithm='SAMME.R', estimator=None, learning_rate=1.0,
                    n_estimators=50, random_state=123)
```

```
[216]: s.evaluate_model(best)
```

```
interactive(children=(ToggleButtons(description='Plot Type:', icons=('',),
options= (('Pipeline Plot', 'pipelin...
```

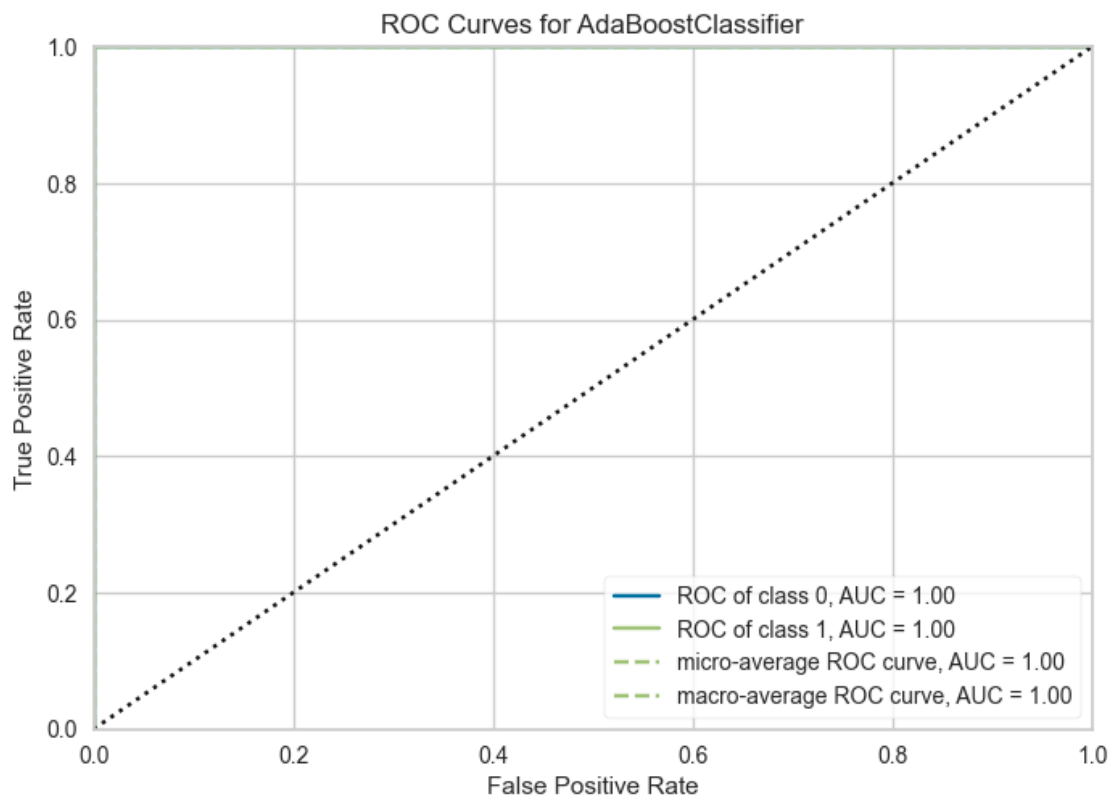

```
[217]: # Save specific plots using plot_model with save=True


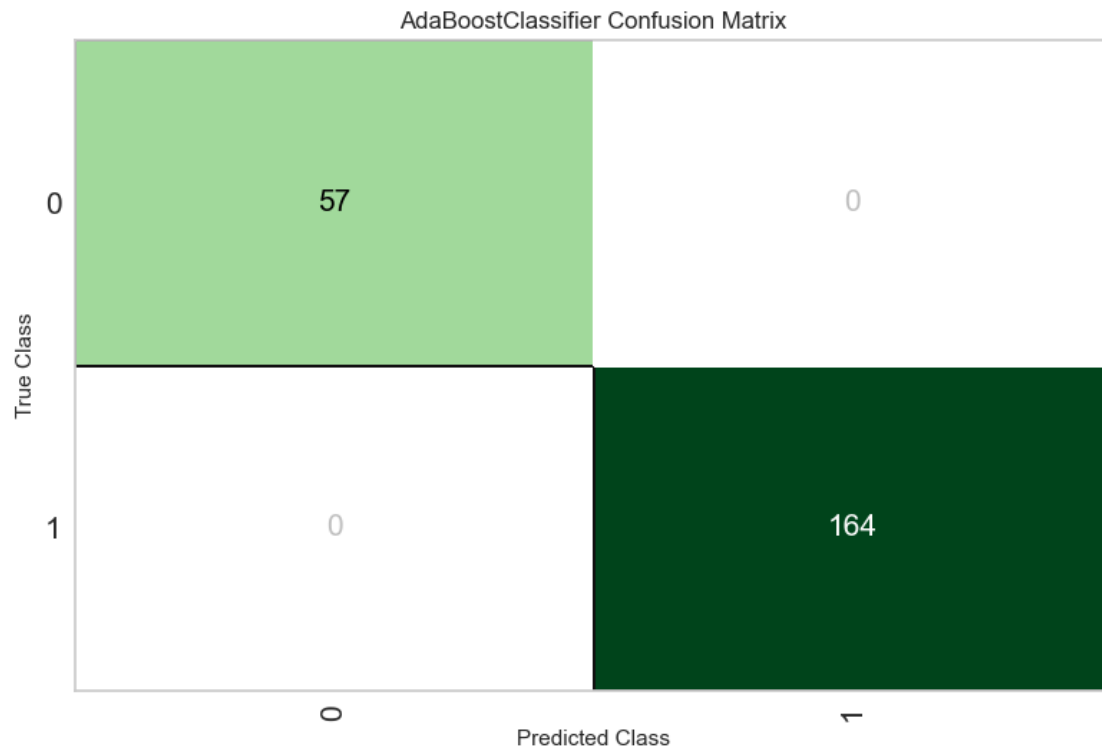
# List of plot types to include
plot_types = ['auc', 'confusion_matrix', 'feature', 'learning']

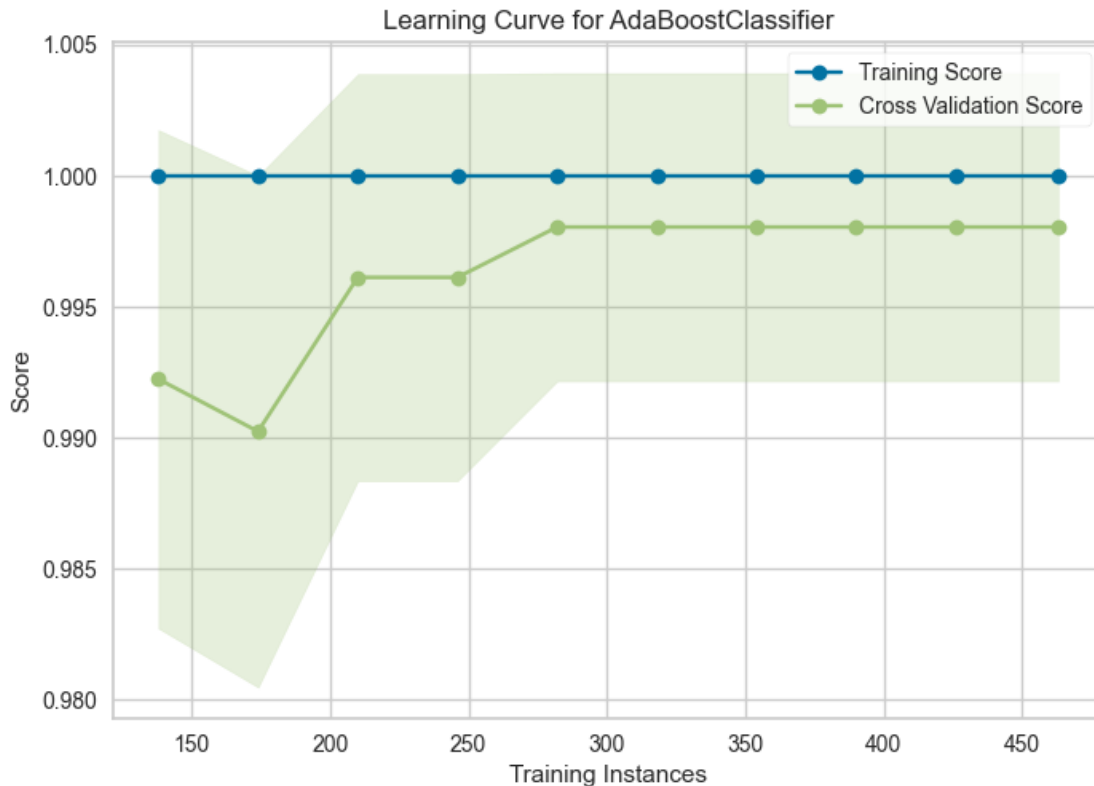
# Iterate over plot types and save each one
for plot_type in plot_types:
    plot_model(best, plot=plot_type, save=True, verbose=False)
```

```
[218]: #DISPLAY THE PLOTS
#Import the package
from IPython.display import Image, display

# Display the PNG images
display(Image('Calibration Curve.png'))
display(Image("AUC.png"))
display(Image("Confusion Matrix.png"))
display(Image("Feature Importance.png"))
display(Image("Learning Curve.png"))
```







The Predict Hold DATA (Testing/assess performace) The predict_holdout is a portion of the dataset that is intentionally set aside and not used during the training of the model. This set is reserved for evaluating the model's performance on unseen data.

```
[219]: #predict/testing
predict_holdout =s.predict_model(best) #Making predictions on the holdout set_
        ↳using the best model (predict_model)

#Retrieve the metrics dataframe
metrics_df = pull() #Calculating various classification metrics based on the_
        ↳predictions made
print(metrics_df) #Printing the metrics dataframe to the console
```

<IPython.core.display.HTML object>

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Ada Boost Classifier	1.0	1.0	1.0	1.0	1.0	1.0	1.0

```
[220]: # Predicting on new data

# Creating a copy of the original data without the 'IMDs' column
New_IMD_data = IMD_data.copy().drop('IMDs', axis=1)
```

```
# Making predictions on the new data using the best-performing model
New_Predictions = s.predict_model(best, New_IMD_data)

# Displaying the first 5 rows of the predictions
New_Predictions.head(5)
```

<IPython.core.display.HTML object>

```
[220]:
```

	age	agecatak	AgeCategory	sex1	religion1	childeduc1	\
0	11.668720	9 to 12	Adolescent	Female	Born Again	Pre-primary	
1	15.000000	13 to 17	Adolescent	Male	Protestant	Pre-primary	
2	6.020534	9 to 12	Adolescent	Female	Muslim	Pre-primary	
3	5.242984	5 to 8	Children	Female	Catholic	Pre-primary	
4	11.696100	9 to 12	Adolescent	Female	Muslim	Pre-primary	

	heightst1	weightst1	BMI_category	childtrib1	...	\
0	134.000000	31.0	Underweight	Non-Munganda but Ugandan	...	
1	137.000000	35.0	Normal weight	Non-Munganda but Ugandan	...	
2	120.099998	19.0	Underweight	Muganda	...	
3	99.000000	15.0	Underweight	Non-Munganda but Ugandan	...	
4	131.600006	30.0	Underweight	Muganda	...	

	tlbase	stin2vntr_	httlpr1	HTTLPRrs35531	rs35531	ptsd	gad	mdd	\
0	1.085469	NaN	3.0	NaN	NaN	0.0	0.0	1.0	
1	1.206384	3.0	3.0	7.0	1.0	0.0	4.0	6.0	
2	1.251031	1.0	1.0	2.0	2.0	0.0	0.0	3.0	
3	1.682923	3.0	2.0	6.0	2.0	0.0	0.0	0.0	
4	1.131768	3.0	1.0	1.0	1.0	1.0	3.0	3.0	

	prediction_label	prediction_score
0	Yes	0.6422
1	Yes	0.9123
2	Yes	0.7084
3	No	0.6625
4	Yes	0.9782

[5 rows x 37 columns]